

Runbook

Vander Lindes Conversational AI

This runbook provides detailed instructions on setting up, deploying, and maintaining the **Vander-lindes Conversational AI** project, a multi-agent conversational AI system designed for natural, context-aware, and policy-compliant customer support for the airline industry.

Table of Contents

- Pre-requisites
- Environment Setup
- Running the Application
- Troubleshooting
- Maintenance
- Deployment
- Backup & Restore

Pre-requisites

Before you begin, ensure you have the following:

1. **Python 3.9+**
2. **Git:** To clone the repository.
3. **LLM API Key:** (e.g., from Google AI Studio) for response generation.

Environment Setup

❖ Clone the Repository

Clone the GitHub repository using the following command:

```
git clone https://github.com/Surya-S-17/Vander-lindes-Conversational-AI.git
cd Vander-lindes-Conversational-AI
```

❖ Create a Virtual Environment

It's recommended to set up a Python virtual environment to isolate dependencies:

```
python -m venv venv
source venv/bin/activate # On Windows use `venv\Scripts\activate`
```

❖ Install Dependencies

Install the required Python dependencies listed in

`Requirements.txt`

`pip install -r requirements.txt`

❖ Set Environment Variables

Create a `.env` file in the root directory to store your API keys and other sensitive information:

`LLM_API_KEY=your_api_key_here`

Replace `your_api_key_here` with your actual API key.

Running the Application

The application consists of several backend services and a frontend demo. You will need to run each component in separate terminals.

Backend Services

You need to run each of the four backend services in separate terminals:

- **Intent Classifier Service** (Terminal 1):
`uvicorn intent_classify:app --reload`
- **Context Management Service** (Terminal 2):
`uvicorn context_management:app --reload`
- **Policy Retrieval (RAG) Service** (Terminal 3):
`uvicorn Policy_retrival_RAG:app --reload`
- **Response Generation Service** (Terminal 4):
`uvicorn response_generation:app --reload`

Frontend Demo

Once the backend services are running, you can start the frontend:

- **Launch Streamlit Demo** (Terminal 5):
`streamlit run demo.py`
- **Open the Demo in Browser:**
Navigate to `http://localhost:8501` in your browser to interact with the chatbot.

Troubleshooting

Here are some common issues and their solutions:

Issue 1: **ModuleNotFoundError** for required libraries

- **Solution:** Ensure that all dependencies are installed:
`pip install -r requirements.txt`

Issue 2: Backend services fail to start

- **Solution:** Ensure all backend services are running in separate terminals. If a service fails, check the logs for specific errors (e.g., incorrect configuration, missing dependencies).

Issue 3: Cannot connect to the frontend

- **Solution:** Make sure the backend services are running. The frontend requires all backend services to be active to function properly.

Issue 4: **LLM_API_KEY** not set correctly

- **Solution:** Check the `.env` file for correctness. Ensure the API key is properly assigned:
`LLM_API_KEY=your_api_key_here`

Maintenance

- **Update Dependencies**

If new updates to dependencies are released, run:

```
pip install --upgrade -r requirements.txt
```

- **Monitor Backend Services**

Ensure the backend services are running smoothly. You can set up logging to track errors in services like:

```
intent_classify:app
context_management:app
Policy_retrival_RAG:app
response_generation:app
```

- **Check for Errors in Response Generation**

If the responses from the bot seem off or incorrect, ensure the underlying LLM model is correctly configured and the API key is active.

- **Database & RAG Maintenance**

If using a vector database (e.g., ChromaDB), ensure it is up to date with relevant data for Policy Retrieval. You may need to reindex periodically to improve search and retrieval performance.

Deployment

For deployment to a production environment (e.g., AWS, Azure, Google Cloud), follow these steps:

1. **Containerization with Docker**

You can containerize the application for easier deployment using Docker. Create a `Dockerfile` for each microservice, or use a single `docker-compose.yml` file to orchestrate the containers.

2. **Configure Web Server**

Deploy the backend services on a web server like **Nginx** or **Gunicorn** to handle production traffic. If using **FastAPI**, deploy via `uvicorn` with **Gunicorn** for better performance.

3. **Set up SSL/TLS:**

For secure communication, ensure SSL/TLS is set up on your web server for HTTPS.

4. **Load Balancer:**

If deploying multiple instances for scalability, configure a load balancer to distribute the traffic across instances.

5. **Environment Configuration:**

Set up environment variables in your production environment (e.g., AWS Lambda, Google Cloud Functions, or Heroku).

Backup & Restore

1. **Backup Backend Data:**

If using a database for policy storage (e.g., ChromaDB), schedule regular backups to avoid data loss. You can use cloud backup services or manual backups based on your storage configuration.

2. **Backup Configuration Files:**

Make sure to backup critical configuration files such as `.env`, `requirements.txt`, and any custom scripts that define the logic of your agents.

3. **Restore from Backup:**

To restore from a backup, follow the reverse process to restore the configuration files and data backups to the correct locations.

Conclusion

This runbook will guide you through the setup, troubleshooting, deployment, and maintenance of the **Vander-lindes Conversational AI** project. Follow these steps for a smooth operational experience and refer back to this document whenever necessary for guidance.